

Wheel of Fortune - Theory

Tansel Arif

tanselarif@live.co.uk

February 11, 2019

Introduction

It is sometimes the case that a random variable is dependent upon another random variable. For example, on some slot machines, the number of spins of the bonus wheel depends on the number of spin/bonus icons you achieve on the slots wheel itself. If you get 1 spin icon you get to spin the wheel once, if you get 2 spin icons, you get to spin the wheel twice and so on. If there is a certain probability of winning the bonus in the wheel each time you spin it and the wins from each spin are independent, then the more times you're allowed to spin it, the more chance of winning the bonus. But, the number of initial spin icons determines the number of times you're allowed to spin the wheel which in turn give you more chances to win the bonus. The question is, what is the probability of winning the bonus each time you play?

Additionally, suppose the probability distribution of the slots is known but not that of the wheel. Given observations on the history of the games played and the results of those games, it is possible to make some guess on the probability distribution of the wheel that would most likely result in such data. We will approach this part of the problem by thinking in a Bayesian context.

Below, I frame a similar problem with a wheel and coin flips to simplify the work a little.

1 The Problem

You spin the wheel of fortune. The wheel gives 0 with probability $\frac{1}{20}$, 1 with probability $\frac{1}{2}$, 2 with probability $\frac{1}{4}$, 3 with probability $\frac{3}{20}$ and 4 with probability $\frac{1}{20}$.

Depending on the number you get from the wheel, you are allowed to flip a coin that many times and record the number of heads you get. The coin is biased towards heads with a probability of $\frac{7}{10}$ of getting a heads.

What is the probability distribution of the number of heads if you play the game?

2 The Solution

Let W be a random variable (RV) representing the number obtained from spinning the wheel. Its distribution is shown in Figure 1.

Let X_i be independent and identically distributed (i.i.d) RVs representing the outcome of a flip of the coin. Then $X_i \sim \text{Ber}(p)$ where $p = 0.7$ (i.e. each X_i is a Bernoulli random variable with probability of success p).

Define a RV, Y , as

$$Y = \sum_{i=1}^W X_i$$




Figure 1: This is the probability mass function for W .

where W itself is random. We can approach this problem using probability generating functions (p.g.f):

Definition 2.1 *Probability Generating Function* If X is a discrete RV taking values in the non-negative integers $\{0, 1, 2, 3, \dots\}$, then the probability generating function of X is defined as [1]

$$G_X(s) = E(s^X) = \sum_{k=0}^{\infty} s^k P(X = k)$$

Moreover,

Theorem 2.1 (Uniqueness of p.g.f) *The distribution of a RV is uniquely determined by its p.g.f.*

So let's find the p.g.f of W . From Definition 2.1:

$$\begin{aligned}
 G_W(s) &= E(s^W) \\
 &= \sum_{k=0}^4 s^k P(W = k) \\
 &= s^0 P(W = 0) + s^1 P(W = 1) + s^2 P(W = 2) + s^3 P(W = 3) + s^4 P(W = 4) \\
 &= \frac{1}{20} + \frac{1}{2}s + \frac{1}{4}s^2 + \frac{3}{20}s^3 + \frac{1}{20}s^4
 \end{aligned} \tag{1}$$

The p.g.f of $X \sim \text{Ber}(p)$ is

$$\begin{aligned}
 G_X(s) &= E(s^X) \\
 &= \sum_{k=0}^1 s^k P(X = k) \\
 &= s^0 P(X = 0) + s^1 P(X = 1) \\
 &= 1 - p + sp
 \end{aligned} \tag{2}$$

and the p.g.f of $Y = \sum_{i=1}^W X_i$ is

$$\begin{aligned}
G_Y(s) &= E(s^Y) \\
&= E(s^{\sum_{i=1}^W X_i}) \\
&= \sum_{w=0}^4 E(s^{\sum_{i=1}^w X_i} | W = w) P(W = w) \\
&= \sum_{w=0}^4 E(s^{\sum_{i=1}^w X_i}) P(W = w) \\
&= \sum_{w=0}^4 E(\Pi_{i=1}^w s^{X_i}) P(W = w) \\
&= \sum_{w=0}^4 \Pi_{i=1}^w E(s^{X_i}) P(W = w) \\
&= \sum_{w=0}^4 \Pi_{i=1}^w G_X(s) P(W = w) \\
&= \sum_{w=0}^4 G_X(s)^w P(W = w) \\
&= E(G_X(s)^W) \\
&= G_W(G_X(s))
\end{aligned} \tag{3}$$

where the third line is the law of total probability, the fourth line is due to the RVs X_i not depending on W and the sixth line comes from the independence of the X_i s. Using equations 1 and 2 and collecting powers of s

$$\begin{aligned}
G_Y(s) &= \frac{1}{20} + \frac{1}{2}(1-p) + \frac{1}{4}(1-p)^2 + \frac{3}{20}(1-p)^3 + \frac{1}{20}(1-p)^4 \\
&\quad + sp\left(\frac{1}{2} + \frac{1}{2}(1-p) + \frac{9}{20}(1-p)^2 + \frac{1}{5}(1-p)^3\right) \\
&\quad + s^2p^2\left(\frac{1}{4} + \frac{9}{20}(1-p) + \frac{3}{10}(1-p)^2\right) \\
&\quad + s^3p^3\left(\frac{3}{20} + \frac{1}{5}(1-p)\right) \\
&\quad + s^4p^4\frac{1}{20}
\end{aligned} \tag{4}$$

Keeping Definition 2.1 in mind, the coefficient of s^i in Equation 4 is $P(Y = i)$. The probability distribution of Y is then

$$\begin{aligned}
P(Y = 0) &= \frac{1}{20} + \frac{1}{2}(1-p) + \frac{1}{4}(1-p)^2 + \frac{3}{20}(1-p)^3 + \frac{1}{20}(1-p)^4 \\
P(Y = 1) &= p\left(\frac{1}{2} + \frac{1}{2}(1-p) + \frac{9}{20}(1-p)^2 + \frac{1}{5}(1-p)^3\right) \\
P(Y = 2) &= p^2\left(\frac{1}{4} + \frac{9}{20}(1-p) + \frac{3}{10}(1-p)^2\right) \\
P(Y = 3) &= p^3\left(\frac{3}{20} + \frac{1}{5}(1-p)\right) \\
P(Y = 4) &= p^4\frac{1}{20}
\end{aligned} \tag{5}$$

Notice that if the probability of heads in the coin flip is 1 ($p = 1$), we recover the probability distribution of W , i.e. W completely determines the number of heads.

Conversely, if the probability of heads in the coin flip is 0 ($p = 0$), the coin will always land tails regardless of the outcome of a trial from W . Then $P(Y = 0) = 1$, i.e. no matter what happens, there will be no heads from the game.

So, to answer the original question, we plug $p = 0.7$ into Equation 5 to obtain the following distribution for the number of heads:

$$\begin{aligned} P(Y = 0) &= 0.226955 \\ P(Y = 1) &= 0.48713 \\ P(Y = 2) &= 0.20188 \\ P(Y = 3) &= 0.07203 \\ P(Y = 4) &= 0.012005 \end{aligned} \tag{6}$$

3 Finding the probability p from observations

Suppose that we are not told the value of p . Given an observation of W , Y has a more well-known distribution

$$Y = \sum_{i=0}^w X_i$$

In particular, Y is just the sum of independent Bernoulli ($Ber(p)$) RVs which has a binomial distribution

$$Y \sim B(w, p)$$

(See section A for proof using p.g.f definitions above). In mathematical notation, the above means

$$Y|W = w \sim B(w, p)$$

Let $p \sim Beta(\alpha, \beta)$ for some α, β . I have chosen this distribution for p as this is conjugate to the Binomial Distribution (see section B) which makes calculations easier.

The posterior distribution for p can then be written as

$$\begin{aligned} P(p|Y, W) &= \frac{P(p, Y, W)}{P(Y, W)} \\ &= \frac{P(Y|W, p)P(W)P(p)}{\sum_{w=0}^4 P(Y|W = w)P(W = w)} \\ &= Beta(p, \alpha + \sum_{i=1}^n y_i, \beta + \sum_{i=1}^n w_i - \sum_{i=1}^n y_i) \end{aligned} \tag{7}$$

where in the last line we used the conjugacy property (see section B). If we were just interested in finding the most likely value for p , we wouldn't have needed to restrict ourselves to conjugate priors since the calculation of the denominator in the second line of Equation 7 would not be required; the main advantage of using conjugate priors in non-Monte Carlo scenarios. This would allow a larger variety of prior distributions for p .

Let us find the most likely distribution for p . In order to find p which maximises $P(p|Y, W)$, we need only concern ourselves with the terms involving p . Using the definition of the *Beta* distribution, we can write Equation 7 as

$$P(p|Y, W) = K p^{\alpha'-1} (1-p)^{\beta'-1} \tag{8}$$

where K is a constant and

$$\alpha' = \alpha + \sum_{i=1}^N y_i, \beta' = \beta + \sum_{i=1}^N w_i - \sum_{i=1}^N y_i$$

Taking the derivative of Equation 8 with respect to p

$$\frac{\partial P(p|Y, W)}{\partial p} = \left[(\alpha' - 1)p^{\alpha'-2}(1-p)^{\beta'-1} - p^{\alpha'-1}(\beta' - 1)(1-p)^{\beta'-2} \right] \quad (9)$$

Replacing α' and β' , and setting this equal to zero we obtain

$$\begin{aligned} p &= \frac{1 - \alpha'}{2 - \alpha' - \beta'} \\ &= \frac{1 - (\alpha + \sum_{i=1}^N y_i)}{2 - (\alpha + \sum_{i=1}^N y_i) - (\beta - \sum_{i=1}^N w_i - \sum_{i=1}^N y_i)} \\ &= \frac{1 - (\alpha + \sum_{i=1}^N y_i)}{2 - \alpha - \beta - \sum_{i=1}^N w_i} \end{aligned} \quad (10)$$

Equation 10 gives us a way of calculating the value of p that is most likely given observations of W and Y . Going further, we may only have data for the overall outcomes of the trials, i.e. the RV Y . In this case we may utilise the MLE of the mean of outcomes of W to be used in Equation 10 as follows

$$\begin{aligned} p &= \frac{1 - (\alpha + \sum_{i=1}^N y_i)}{2 - \alpha - \beta - NE[W]} \\ &\approx \frac{1 - (\alpha + \sum_{i=1}^N y_i)}{2 - \alpha - \beta - N \frac{\sum_{i=1}^N w_i}{N}} \end{aligned} \quad (11)$$

Let's see what happens when we set $p = 2$ in the theory section above and provide 'hypothetically ideal' observations to be used in Equation 11. The probability distribution of Y then becomes

$$\begin{aligned} P(Y = 0) &= 0.70728 \\ P(Y = 1) &= 0.25808 \\ P(Y = 2) &= 0.03208 \\ P(Y = 3) &= 0.00248 \\ P(Y = 4) &= 0.00008 \end{aligned} \quad (12)$$

According to 12, out of $N = 100000$ trials, we expect to have: 70728 trials resulting in 0 heads, 25808 trials resulting in 1 head, 3208 trials resulting in 2 heads, 248 trials resulting in 3 heads and 8 trials resulting in 4 heads for a total of 33000 heads in 100000 trials/games. Additionally,

$$E[W] = 0 * \frac{1}{20} + 1 * \frac{1}{2} + 2 * \frac{1}{4} + 3 * \frac{3}{20} + 4 * \frac{1}{20} = 1.65$$

For the prior distribution for p , we may specify the parameters $\alpha = 1$ and $\beta = 1$ resulting in a *Beta* distribution that is uniform without assuming any information about the unbiased/bias nature of the coin. The prior distribution for p is shown in blue in Figure 2. Using these parameters, our estimation for the value of p that was used to generate these observations is

$$p = \frac{1 - (1 + 33000)}{-100000 \times 1.65} = 0.2$$

Note that here we have only used the knowledge of the outcomes along with the p.d.f of W to determine p . We did not need to know the distribution of Y . In general, the p.d.f of W may be unknown requiring the data from the outcomes of W , or there may be strong prior evidence that the coin is unbiased (we could set $\alpha = 50, \beta = 50$ for example to assume a strong prior around $p = 0.5$ (see Figure 2)). Additionally, a posterior probability distribution for p may be required, in which case we have $p \sim \text{Beta}(1 + 33000, 1 + 1.65 * 100000 - 33000)$ shown in Figure 3.




Figure 2: This is the probability distribution of $Beta(1, 1)$ and $Beta(50, 50)$.




Figure 3: This is the posterior probability distribution for p .

4 Simulations

We can perform simulations of this process by first selecting a value, w , from the distribution for W then selecting w values from a Bernoulli distribution with specified probability of success p .

Below shows Python code used for performing the simulation:

```
# Import modules
from scipy import stats
import math
import numpy

# Set the seed
np.random.seed(101)

# Build W by concatenation of lists
w = []
w.append(0)
w.extend([1]*10+[2]*5+[3]*3+[4]*1)

# Set the probability of heads from a coin flip
p = 0.2

# Set the number of games
n = 10000

# Initialise the total number of heads obtained
y_outcome = 0

# Get n samples from w
w_outcome = np.random.choice(w,size=n)

# For each w, flip the coin that many times
y_outcome = [np.sum(stats.bernoulli.rvs(p,size=this_w)) for this_w in w_outcome]

y_sum = sum(y_outcome)
w_sum = sum(w_outcome)

print('Total number of heads = {}'.format(y_sum))
print('Total number of coin flips = {}'.format(w_sum))

>>> Total number of heads = 3202
>>> Total number of coin flips = 16536
```

Next we plot a representation of the posterior distribution for p using the theory in the previous section.

```
# Get figure object
fig = plt.figure(figsize=(5, 5))

# Get axes object
axes = fig.add_axes([0.2,0.2,0.8,0.8])

# Create a numpy array of 100 points equally spaced
x = np.linspace(0,1,100)

# Create a numpy array of the posterior distribution using actual observed values of W
y = np.array([stats.beta.pdf(i,1+y_sum,1-y_sum+w_sum) for i in x])

# Create a numpy array of the posterior distribution by estimating sum(W) with it's expected value
z = np.array([stats.beta.pdf(i,1+y_sum,1-y_sum+n*1.65) for i in x])
```

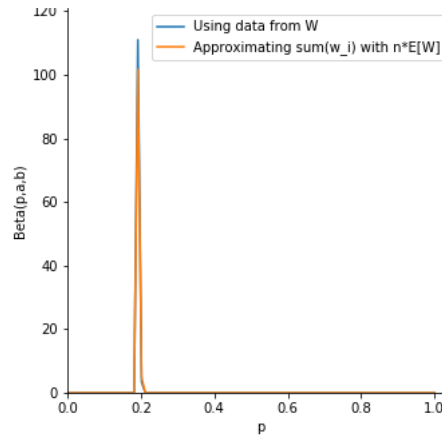


Figure 4: A simulation is run by performing $n = 10000$ trials of the process of spinning a wheel and flipping a coin a number of times specified by the value obtained from the wheel and summing the number of heads obtained throughout the simulation. The probability of heads is chosen as $p = 0.2$. The theory resulting in Equation 7 is used to generate both probability distributions in the plot. The blue line uses the data from the actual trials from W whereas the orange line uses an estimated value for W for cases where we only know the outcome Y .

```
# Plot onto the axes
axes.plot(x, y, '-o',ms=0,label='Using data from W')
axes.plot(x, z, '-o',ms=0,label='Approximating sum(w_i) with n*E[W]')

# Set the axis labels
axes.set_xlabel('p')
axes.set_ylabel('Beta(p,a,b)')

# Show legend
axes.legend()

# Set axis limits
axes.set_xlim(0,1.05)
axes.set_ylim(0,max(y.max(),z.max())+10)

# Remove some borders
axes.spines['right'].set_visible(False)
axes.spines['top'].set_visible(False)
```

This produces the plots in Figure 4. It can be seen that the posterior distribution correctly estimates the probability associated with getting a heads from the coin flip p as 0.2.

A The sum of n Bernoulli random variables is a Binomial random variable

By Equation 2, the p.g.f of a Bernoulli RV X is

$$G_X(s) = 1 - p + sp$$

Similarly, using Definition 2.1, the p.g.f of a Binomial RV Z such that $Z \sim B(n, p)$ and $P(Z = z; n, p) = \binom{n}{z} p^z (1 - p)^{n-z}$ is

$$\begin{aligned} G_Z(s) &= E[s^Z] \\ &= \sum_{z=0}^n s^z P(Z = z) \\ &= \sum_{z=0}^n s^z \binom{n}{z} p^z (1 - p)^{n-z} \\ &= \sum_{z=0}^n \binom{n}{z} (sp)^z (1 - p)^{n-z} \\ &= (sp + 1 - p)^n \end{aligned} \tag{13}$$

where the last equality is from the identity $(a + b)^n = \sum_{i=0}^n \binom{n}{i} a^i b^{n-i}$. Let $Y = \sum_{i=0}^n X_i$, i.e. the RV Y is the sum of the n RVs X_i . Then the p.g.f of Y is

$$\begin{aligned} G_Y(s) &= E[s^Y] \\ &= \sum_{y=0}^n s^y P(Y = y) \\ &= E[s^{\sum_{i=0}^n X_i}] \\ &= E[\prod_{i=0}^n s^{X_i}] \\ &= \prod_{i=0}^n E[s^{X_i}] \quad (\text{since } X_i \text{ are independent}) \\ &= (G_X(s))^n \quad (\text{by the p.g.f of a Bernoulli RV}) \\ &= (sp + 1 - p)^n \end{aligned} \tag{14}$$

This p.g.f is the same as that of the Binomial RV above and by Theorem 2.1, $Y \sim B(n, p)$.

B Conjugacy property of $Y \sim B(n, p)$ and $p \sim \text{Beta}(\alpha, \beta)$

Let W be a RV with a defined distribution with parameters specified. Let $Y \sim B(n, p)$ and prior $p \sim \text{Beta}(\alpha, \beta)$. Then

$$\begin{aligned}
P(p|Y, W) &= \frac{P(p, Y, W)}{P(Y, W)} \\
&= \frac{P(Y|W, p)P(W)P(p)}{\sum_{w=0}^4 P(Y|W = w)P(W = w)} \\
&= \frac{P(Y_1 = y_1, Y_2 = y_2, \dots, Y_n = y_n|W, p)P(W)P(p)}{\sum_{w=0}^4 P(Y|W = w)P(W = w)} \\
&= \frac{P(Y_1 = y_1|W, p) \dots P(Y_n = y_n|W, p)P(W)P(p)}{\sum_{w=0}^4 P(Y|W = w)P(W = w)} \quad \text{since the } Y_i \text{ are independent} \\
&= \frac{P(Y_1 = y_1|W_1 = w_1, p) \dots P(Y_n = y_n|W_n = w_n, p)P(W)P(p)}{\sum_{w=0}^4 P(Y|W = w)P(W = w)} \quad \text{since } Y_i \text{ depends only on } W_i \\
&= \frac{B(y_1; w_1, p) \dots B(y_n; w_n, p)P(W)Beta(p, \alpha, \beta)}{\sum_{w=0}^4 P(Y|W = w)P(W = w)} \\
&= \frac{\binom{w_1}{y_1} p^{y_1} (1-p)^{w_1-y_1} \dots \binom{w_n}{y_n} p^{y_n} (1-p)^{w_n-y_n} P(W) \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} p^{\alpha-1} (1-p)^{\beta-1}}{\sum_{w=0}^4 P(Y|W = w)P(W = w)} \\
&= \frac{\frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} P(W) \prod_{i=0}^n \left[\binom{w_i}{y_i} \right]}{\sum_{w=0}^4 P(Y|W = w)P(W = w)} p^{\sum_{i=0}^n y_i + \alpha - 1} (1-p)^{\sum_{i=0}^n (w_i - y_i) + \beta - 1} \\
&= K p^{\sum_{i=0}^n y_i + \alpha - 1} (1-p)^{\sum_{i=0}^n (w_i - y_i) + \beta - 1}
\end{aligned} \tag{15}$$

The K is a constant and not dependent upon p . Here, we note that the posterior distribution $P(p|Y, W)$ is itself a probability distribution and must sum/integrate to 1. Since Equation 15 has the form of a $Beta(p, \alpha + \sum_{i=1}^n y_i, \beta + \sum_{i=1}^n w_i - \sum_{i=1}^n y_i)$ distribution and must integrate to 1, we have that $p|Y, W \sim Beta(\alpha + \sum_{i=1}^n y_i, \beta + \sum_{i=1}^n w_i - \sum_{i=1}^n y_i)$ and K is a proportionality parameter which serves to scale the expression so that it has unit area. Here we have seen how a prior distribution for p has been updated with data from the distribution for Y resulting in a posterior distribution which remains within the same family as the prior distribution. The *Binomial* and *Beta* distributions thus have a property called Conjugacy.

References

- [1] Wikipedia. Probability-generating function. https://en.wikipedia.org/wiki/Probability-generating_function.